

- RAG combines retrieval of relevant documents with generative models to improve response quality.
- It uses a two-step process: first retrieving information, then generating a coherent answer.
- The retrieval component fetches contextually relevant data from a large corpus.
- The generative model synthesizes this information into a natural language response.
- RAG is particularly useful for tasks requiring up-to-date or specific knowledge.
- It helps mitigate the limitations of generative models by grounding responses in real data.